VOL. 2

PALANTIR x FCA / AI-AUGMENTED SUPERVISION SERIES

ETHICAL COMPROMISES

The Moral Architecture of
AI-Augmented Financial Supervision

FROM THE POSITION OF A UK FCA-REGULATED PSP
HIGH-RISK MERCHANT ACQUIRING · LIMITED COMPLIANCE RESOURCES

ETHICAL ANALYSIS March 2026

7	**£500M+**	**42,000**
Ethical Dimensions Examined	Palantir UK Public Contracts	Firms' Data Exposed

CONTENTS

STRUCTURE OF VOLUME 2

1 EXECUTIVE SUMMARY

SEVEN ETHICAL FAULT LINES IN AI-AUGMENTED SUPERVISION

Volume 1 of this series analysed the consequences of the FCA's Palantir deployment. Volume 3 traces the structural preconditions that made it inevitable. This volume occupies the space between: the ethical compromises embedded in the decision itself and in the supervisory architecture it creates.

These are not abstract philosophical concerns. For a small PSP with limited resources, each ethical fault line translates into a concrete operational dilemma — a situation where doing the right thing is structurally ambiguous, where compliance and justice may point in different directions, and where the firm's survival depends on navigating a moral landscape it did not design and cannot reshape.

VENDOR ETHICS

The FCA has entrusted the UK's financial intelligence to a company accused by Amnesty International, the UN Special Rapporteur, and its own former employees of enabling human rights abuses. Does the ethical provenance of the analytical tool contaminate the legitimacy of supervisory outcomes it produces?

ALGORITHMIC JUSTICE

AI-driven detection inverts the traditional sequence of suspicion. The algorithm identifies statistical anomalies first; human judgment follows. For the flagged firm, the practical experience is of being treated as guilty until it can demonstrate innocence — at its own cost.

DATA SOVEREIGNTY

A US-incorporated defence contractor, co-founded with CIA seed funding, now processes the UK financial regulator's most sensitive intelligence. The FCA insists on UK data residency and processor-only status. But sovereignty is not merely a question of where data is stored.

THE ACCOUNTABILITY VOID

When an AI system flags a firm, triggers a Section 166 review costing £100K–500K+, and the flag proves false — who bears the cost? The current architecture provides no mechanism for redress against algorithmic misidentification.

PROPORTIONALITY AS MORAL PRINCIPLE

Applying identical compliance standards to a 5-person PSP and a multinational bank is formally equitable and substantively unjust. AI supervision intensifies this injustice by making previously invisible gaps visible.

OUTSOURCED JUDGMENT

When a commercial vendor's platform shapes what the regulator can see and how anomalies are defined, the boundary between public regulatory authority and private commercial interest blurs in ways that have no clear precedent in UK administrative law.

THE IMPOSSIBLE POSITION

Small firms must simultaneously comply with a regime designed for institutional scale, accept surveillance by a vendor they cannot vet, bear the cost of algorithmic error, and maintain the pretence that this is a level playing field.

2 THE VENDOR ETHICS PROBLEM

PALANTIR'S HUMAN RIGHTS RECORD AND THE LEGITIMACY QUESTION

The question is not whether Palantir builds effective analytical software. It manifestly does. The question is whether the ethical provenance of the vendor materially affects the legitimacy of the supervisory outcomes its technology produces — and whether regulated firms have standing to raise this concern.

The Documented Record

US IMMIGRATION (ICE)

Amnesty International has determined there is a "high risk that Palantir is contributing to serious human rights violations of migrants and asylum seekers." In 2017, ICE relied on Palantir technology to arrest parents and caregivers of unaccompanied children. In 2025, Palantir's ImmigrationOS was identified by Amnesty as a key component of the "Catch and Revoke" programme, automating visa revocations with minimal due process.

MILITARY / ISRAEL-PALESTINE

The UN Special Rapporteur on Palestinian territories named Palantir as an enabler of "unlawful use of force," alleging real-time battlefield AI systems and predictive policing tools supplied to the Israeli military. Amnesty International identified Palantir among companies enabling or profiteering from what it characterised as genocide, occupation, and apartheid.

PREDICTIVE POLICING

Palantir's Gotham platform has been deployed in US cities for predictive policing. After a six-year secret contract was exposed in New Orleans, the city cancelled the programme. Research indicates these technologies entrench systemic bias and disproportionately affect minority communities.

NHS DATA CONCERNS

The £330M Federated Data Platform contract has drawn sustained criticism from health professionals and human rights groups. Medact's 2026 briefing argues that NHS engagement with Palantir "may contravene ethical procurement practices" and risk linking NHS procurement to alleged human rights abuses.

INTERNAL DISSENT

In May 2025, thirteen former Palantir employees publicly condemned the company, stating that its founding principles "have now been violated." They warned that "government databases are already erasing references to transgender people" and that "these injustices could be facilitated by the very software infrastructure we helped build."

The Legitimacy Paradox

The FCA's mandate is to protect consumers, ensure market integrity, and promote competition. It exercises coercive powers — the ability to fine, restrict, and de-authorise firms — that derive their legitimacy from public trust in the regulator's integrity and processes. When those powers are mediated through technology supplied by a vendor whose own ethical conduct is subject to serious, documented allegations from the UN, Amnesty International, and the vendor's own former employees, a legitimacy question arises that no data residency clause can resolve.

This is not to suggest that the FCA's supervisory outcomes are necessarily tainted. But it is to observe that the choice of vendor is not ethically neutral, and that regulated firms — particularly those subjected to adverse supervisory action informed by Palantir's analytics — may have grounds to question whether the process meets the standards of fairness that UK administrative law requires.

3 ALGORITHMIC JUSTICE

PRESUMPTION OF INNOCENCE IN THE AGE OF PATTERN DETECTION

The traditional model of financial supervision follows a broadly judicial sequence: suspicion arises from specific intelligence, investigation is initiated, evidence is gathered, and conclusions are reached. The firm is presumed compliant until demonstrated otherwise. AI-driven supervision inverts this sequence in ways that have profound implications for procedural justice.

The Inversion Mechanism

In the AI model, the algorithm processes the entire data lake and identifies statistical deviations from established baselines. A firm is flagged not because specific intelligence suggests wrongdoing, but because its data profile is anomalous relative to its peer group. The investigation follows the flag, not the suspicion. The firm learns it is under scrutiny — perhaps through an information request, a Section 166 notice, or a Voluntary Requirement — without being told what specific conduct triggered the attention, because there may be no specific conduct: only a pattern that deviates from the norm.

The Practical Impact on Small PSPs

BURDEN INVERSION

The firm must now demonstrate that its compliance is adequate — a reversal of the traditional burden where the regulator demonstrates inadequacy. For a small PSP with limited resources, assembling evidence of compliance in response to an algorithmically generated inquiry can consume months of senior management time.

OPAQUE TRIGGERS

The firm may never learn precisely what algorithmic signal triggered the regulatory interest. Palantir Foundry's cross-referencing produces composite risk signals from multiple data sources. The firm sees the investigation; it does not see the algorithm's reasoning.

PEER BENCHMARKING FALLACY

AI anomaly detection relies on establishing baselines from peer groups. But "peers" are defined by the algorithm's categorisation. A small PSP in high-risk acquiring may be benchmarked against firms with fundamentally different risk profiles, transaction volumes, or business models — producing false deviations that trigger genuine investigations.

THE COST OF INNOCENCE

A Section 166 review costs the firm £100K–500K+ regardless of outcome. There is no mechanism for recovering costs if the algorithmic flag proves unfounded. The firm pays the price of the algorithm's false positive from its own operating capital.

THE ETHICAL CORE

Traditional AML systems produce false positive rates of 90–95%. AI reduces this but does not eliminate it. The ethical question is not whether false positives occur — they inevitably will — but who bears the cost when they do. In the current architecture, the cost falls entirely on the firm. The regulator and the vendor bear no financial consequence for algorithmic error. This is not a technical limitation; it is an ethical choice embedded in the system design.

4 DATA SOVEREIGNTY

A US DEFENCE CONTRACTOR PROCESSING UK FINANCIAL INTELLIGENCE

The FCA has structured the Palantir contract with explicit data protection safeguards: processor-only status, UK data residency, FCA-owned encryption keys, mandatory data destruction at contract end. These are necessary conditions. The question is whether they are sufficient.

The Structural Concern

Palantir Technologies Inc. is incorporated in the United States. It was co-founded with seed capital from In-Q-Tel, the venture capital arm of the CIA. Its CEO has publicly described the company's mission in terms of Western military and intelligence supremacy. Its largest contracts are with the US Department of Defense, ICE, and the intelligence community. It now processes the UK financial regulator's most sensitive intelligence — case files on problem firms, suspected fraud data, SARs, investigation emails, and call recordings across 42,000 regulated entities.

Data sovereignty is not merely a question of physical storage location. It encompasses the jurisdictional reach of foreign governments over the vendor, the vendor's existing obligations to foreign intelligence services, the institutional culture shaped by intelligence community origins, and the risk that analytical methodologies — even if not the data itself — can be transferred across jurisdictions.

Three Dimensions of Sovereignty Risk

JURISDICTIONAL EXPOSURE

Palantir is subject to US law, including potential national security letters, FISA court orders, and executive directives. While UK data residency provides some protection, the legal boundary between data residency and jurisdictional reach is neither absolute nor untested. The US CLOUD Act of 2018 explicitly asserts US government authority to compel US-incorporated companies to produce data regardless of where it is stored.

ANALYTICAL SOVEREIGNTY

Even if no raw data leaves the UK, the analytical models, ontological structures, and pattern recognition methodologies developed during the trial represent transferable intelligence. Palantir's intellectual property in analytical methods crosses borders even if the data does not. The FCA retains IP from this specific engagement, but Palantir's institutional learning cannot be contractually erased.

CULTURAL SOVEREIGNTY

The ontological models that determine how entities are classified, how relationships are mapped, and how anomalies are defined reflect the institutional assumptions of their creators. A platform designed for US military intelligence and immigration enforcement may embed assumptions about risk, deviance, and threat that do not translate cleanly to UK financial regulation — and those assumptions are not visible in the outputs.

5 THE ACCOUNTABILITY VOID

WHO ANSWERS WHEN THE ALGORITHM IS WRONG?

In a traditional supervisory model, accountability is relatively clear. A human supervisor makes a decision, that decision can be challenged through the Upper Tribunal, and the FCA can be held to account for procedural failures. AI-augmented supervision introduces a structural ambiguity that existing accountability mechanisms are not designed to address.

The Three-Party Problem

In the Palantir-FCA architecture, three parties are involved in any supervisory decision: **Palantir** builds and maintains the analytical platform; the **FCA** configures the parameters and acts on the outputs; and the **regulated firm** bears the consequences. When the system produces an incorrect output — a false positive that triggers investigation of a compliant firm — accountability diffuses across the three parties in ways that leave the harmed firm without effective recourse.

PALANTIR'S POSITION

"We are a data processor, not a data controller. The organisations that license our products define what can and cannot be done with their data." This framing — explicitly stated in Palantir's 2020 blog post and echoed in the FCA contract — positions the vendor as a neutral tool with no responsibility for outcomes.

THE FCA'S POSITION

AI outputs inform but do not determine supervisory decisions. Human reviewers assess the algorithmic flags before action is taken. The regulator retains accountability for the final decision.

THE REALITY

When an algorithm trained on two years of intelligence data across 42,000 firms flags a specific firm as anomalous, the practical probability that a human reviewer will override the flag is low. The system's output creates a cognitive anchor that subsequent human judgment orbits rather than escapes. The "human in the loop" becomes a formality that legitimises rather than genuinely interrogates the algorithmic output.

The result is what might be called **accountability laundering**: the vendor disclaims responsibility by pointing to the regulator's control; the regulator claims human oversight while relying on algorithmic pre-selection; and the firm bears the cost with no effective mechanism for challenging the system that produced the outcome. The accountability does not disappear — it is redistributed onto the party with the least power to absorb it.

6 PROPORTIONALITY AS MORAL PRINCIPLE

WHEN EQUAL STANDARDS PRODUCE UNEQUAL HARM

The FCA's 2025–2030 strategy commits to proportionality: less intensive supervision for firms demonstrably seeking compliance. But proportionality in supervisory intensity does not extend to compliance standards themselves. The substantive requirements for AML systems, sanctions screening, merchant due diligence, safeguarding, and governance are identical regardless of firm size.

The Structural Injustice

Consider two firms: a multinational payment institution with 500 compliance staff, and a small PSP with 5 employees total. Both must maintain robust AML transaction monitoring, comprehensive sanctions screening, documented merchant due diligence, daily safeguarding reconciliations, Consumer Duty evidence packs, and operational resilience frameworks. The cost of compliance as a percentage of revenue is orders of magnitude higher for the small firm. AI-augmented supervision intensifies this disparity by making previously invisible gaps visible — gaps that existed at the small firm not from intent but from structural incapacity.

The AI Amplification Effect

- AI peer benchmarking compares firms within portfolio categories without weighting for resource availability. A small PSP's SAR filing rate is compared against the portfolio average, which is dominated by larger firms with dedicated compliance teams.
- Cross-referencing capabilities expose documentation gaps that manual supervision would never have found — not because the gaps don't matter, but because proportionate manual supervision correctly deprioritises them.
- The cost of responding to AI-triggered inquiries falls disproportionately on smaller firms. A request for information that takes a large firm's compliance team a day to compile may consume a small firm's entire management capacity for weeks.
- Section 166 reviews, triggered at higher rates by data-driven supervision, cost £100K–500K+ regardless of firm size — a rounding error for a major institution, a potential extinction event for a small PSP.

THE MORAL QUESTION

Formal equality of treatment can constitute substantive injustice when the parties being treated equally occupy radically different structural positions. This is not a novel ethical insight — it is foundational to concepts of equity in law, philosophy, and public administration. Yet the architecture of AI-augmented supervision embeds formal equality as a design principle without accounting for its distributive consequences.

7 THE OUTSOURCING OF REGULATORY JUDGMENT

COMMERCIAL INCENTIVES IN PUBLIC OVERSIGHT

Financial regulation is an exercise of sovereign authority. The FCA's powers to investigate, restrict, fine, and de-authorise firms derive from statute and are exercised in the public interest. When the analytical infrastructure that informs these powers is provided by a commercial vendor, a boundary is crossed that raises questions without easy answers.

The Incentive Misalignment

Palantir is a publicly traded company with a market capitalisation of approximately $300 billion and a well-documented strategy of expanding its presence within government institutions. Its commercial interest in the FCA engagement is not merely the £30K+ per week trial fee — it is the conversion of a short-term trial into a long-term, high-value contract. This conversion is more likely if the trial demonstrates high value, which in the context of regulatory supervision means: finding things. A system that finds nothing is unlikely to be renewed.

This does not require any conscious bias from Palantir's engineers. It is a structural incentive embedded in the commercial relationship. The platform's configuration, the sensitivity of anomaly detection thresholds, the breadth of cross-referencing — all of these technical parameters have direct implications for how many firms are flagged. A vendor whose contract renewal depends on demonstrating analytical value has a structural incentive toward sensitivity over specificity: more flags are better than fewer, from the vendor's commercial perspective.

The Democratic Deficit

Parliamentary scrutiny of the FCA's supervisory decisions is already limited. When those decisions are informed by a proprietary analytical platform whose methodology is commercially confidential, the democratic deficit compounds. MPs who questioned the contract — the Liberal Democrats' Daisy Cooper demanding a government investigation, the Greens' Sian Berry calling for it to be blocked — are exercising precisely the scrutiny function that the structure of the arrangement makes difficult to sustain. How do you hold a regulator accountable for decisions informed by an algorithm you cannot inspect?

THE PRECEDENT

The FCA-Palantir contract establishes a precedent for outsourcing the analytical substrate of regulatory judgment to commercial vendors. If this trial converts into a permanent relationship — as Palantir's NHS and MoD engagements have done — the vendor becomes embedded in the regulator's decision-making infrastructure in ways that are practically irreversible. The UK will have created a model in which a foreign-incorporated defence contractor's platform is a structural component of financial regulation.

8 THE SMALL FIRM'S IMPOSSIBLE POSITION

NAVIGATING ETHICS FROM BELOW

A small FCA-regulated PSP engaged in high-risk merchant acquiring faces a set of ethical constraints that are simultaneously real and irresolvable from its structural position.

The Dilemma Matrix

COMPLY OR CHALLENGE?

The firm can invest scarce resources in proactive compliance — accepting the legitimacy of a supervisory regime it did not consent to and cannot influence — or it can challenge the regime's ethical foundations through industry bodies, parliamentary channels, or legal proceedings. The first preserves the firm's survival; the second may protect its principles. Doing both simultaneously is beyond the resources of most small firms.

TRANSPARENCY OR SELF-PROTECTION?

A firm that proactively discloses compliance gaps to the FCA demonstrates good faith — and also hands the regulator evidence that could trigger formal action. A firm that remediates silently protects itself in the short term — but risks more severe consequences if the gap is independently detected by AI analytics.

INVEST IN COMPLIANCE OR INVEST IN BUSINESS?

Every pound spent on RegTech tools, gap analysis, and compliance documentation is a pound not invested in the business that generates the revenue needed to sustain operations. For a firm operating at the margins, the compliance investment needed to survive AI-augmented supervision may itself be the investment that makes the business unviable.

ACCEPT THE VENDOR OR RAISE THE QUESTION?

A firm that raises ethical concerns about Palantir's involvement in FCA supervision risks being perceived as obstructive. A firm that accepts it silently normalises a precedent that may not serve the industry's long-term interests. There is no neutral position.

A Pragmatic Ethics

The analysis in this volume does not argue that a small PSP should refuse to comply, challenge the FCA in tribunal, or publicly denounce Palantir. It argues something more modest but perhaps more useful: that the ethical compromises are real, that acknowledging them honestly is a precondition for navigating them wisely, and that a firm's compliance strategy should be informed by an understanding of the moral architecture — not merely the technical requirements — of the system in which it operates.

A compliance strategy built on the pretence that the system is ethically neutral will be brittle. A compliance strategy that recognises the ethical tensions and designs for them — documenting not just what the firm does but why, maintaining records that demonstrate intent as well as process, engaging with industry bodies on structural reform — will be more resilient precisely because it is more honest.

CONCLUSION

The ethical compromises embedded in the FCA's Palantir deployment are not incidental. They are structural features of a supervisory model that prioritises detection efficiency over procedural equity, formal equality over substantive justice, and institutional convenience over democratic accountability.

None of these compromises are unique to the FCA or to Palantir. They are inherent in any system where AI-driven analytics inform coercive state action against entities with unequal power to contest the outcomes. What makes the FCA-Palantir case distinctive is the scale of the data involved, the severity of the potential consequences for regulated firms, and the specificity of the vendor's ethical record.

THE ETHICAL COMPROMISES ARE NOT RESOLVABLE FROM BELOW

A small PSP cannot fix the accountability void, the proportionality deficit, or the vendor ethics problem. These are systemic features that require systemic responses — legislative reform, regulatory redesign, or judicial intervention.

BUT THEY CAN BE NAVIGATED WITH HONESTY

A firm that understands the ethical terrain it operates in — that recognises the asymmetry between its position and the regulator's, that documents its intent as well as its process, that engages with industry bodies on structural reform — is better positioned than one that simply ticks boxes.

ETHICS IS NOT SEPARATE FROM STRATEGY

In a regime where algorithms cannot distinguish intent from incapacity, the firm's documented ethical reasoning becomes a material asset. Evidence that a compliance gap resulted from resource constraints rather than negligence may not prevent an investigation — but it can shape its outcome.

The firms that navigate this landscape with the greatest integrity will be those that refuse the false comfort of pretending the ethical compromises do not exist — and instead build their compliance on the harder foundation of seeing them clearly.

VOL. 2: ETHICAL COMPROMISES · MARCH 2026

COAGNES SEMANTIC INTELLIGENCE · REGULATORY ANALYSIS DIVISION