

PROMPTING TECHNIQUE GUIDE

Semantic Kernel Prompting

Read complex texts with LLMs — not once, but in layers.

A step-by-step practical guide with GDPR examples.

A common way to approach this is to ask an LLM to 'summarize this regulation' or 'describe/explain this'. It looks complete. But something is missing.

This guide shows a different approach:
build a field of questions first, then read through it.

CoAgnes Framework

March 2026

What happens when you ask "Summarize GDPR Recital 29"?

A TYPICAL PROMPT

Explain GDPR Recital 29.

MODEL RESPONSE (TYPICAL)

Recital 29 of the General Data Protection Regulation (GDPR) focuses on the concept of pseudonymization as a key strategy for enhancing privacy and security. It encourages data controllers to process personal data in a way that the data can no longer be attributed ...

This looks fine. But it leaves important structure implicit:

- Is pseudonymisation mandatory or merely encouraged?
- Does "kept separately" mean logically, physically, or cryptographically?
- Do pseudonymised data remain personal data?
- Who exactly are "authorised persons" and under what constraints?
- Where does this recital end and other provisions begin?

A flat summary collapses the fragment into a single surface. The structure underneath — definitions, boundaries, connections, open zones — disappears.

What if instead...

You could build a field of questions *before* reading, then pass the text through that field — producing not a summary but a **semantic map** showing what the fragment asserts, where it ends, what it connects to, and what it genuinely leaves open?

That's what this protocol does. Let's examine this step by step.

STEP 1**Generate Questions — Not Answers**

Send the fragment with an instruction to produce questions along three axes. Not answers — questions.

A TYPICAL PROMPT

Looking from inside the semantic space of this fragment, identify question vectors directed toward other parts, definitions, or provisions of the document.

Truth

- Questions pointing toward definitions

Deep

- Questions pointing toward boundaries

Connect

- Questions pointing toward the fragment's place in the overall field

GDPR Recital 29: In order to create incentives to apply pseudonymisation when processing personal data, measures of pseudonymisation should, whilst allowing general analysis, be possible within the same controller...

MODEL OUTPUT (EXCERPT)**Truth**

What does GDPR mean by "pseudonymisation"? → Art. 4(5)

Who is a "controller" in this context? → Art. 4(7)

What counts as "additional information"? → Art. 4(5) + Recital 26

...

Deep

Is this an obligation or merely an incentive?

Does "kept separately" mean logical, physical, or cryptographic separation?

Is pseudonymisation alone sufficient for GDPR compliance?

...

Connect

How does this relate to privacy by design (Art. 25)?

How does it connect to security of processing (Art. 32)?

Where is the boundary with anonymisation (Recital 26)? ...

Why questions first? They form the *field of attention* for everything that follows. Truth questions will load context. Deep questions will guide analysis. Connect questions will organize the output.

Review the questions. Add missing ones. Remove noise. This is your moment to shape the reading.

STEP 2

Build Context Through Questions

Now use the Truth questions as vectors to project a structured context — before any analysis.

A TYPICAL PROMPT

Looking from inside the semantic space of the fragment, project the critically necessary components for full semantic understanding, using these questions as anchoring vectors.

Structure

- Truth: critical definitions
- Deep: applicability questions
- Connect: structural relationships
- Service: concrete semantic forms
- Knowledge: related fragments
- Evolution: temporal relationships
- Responsibility: points of incompleteness

--- Questions from Step 1 (Truth section) ---

What does GDPR mean by "pseudonymisation"?

Who is a "controller" in this context?

...

MODEL OUTPUT (CONDENSED EXCERPT)**Truth: definitions**

Pseudonymisation (Art. 4(5)) — processing so data cannot be attributed to a subject without separate additional information. Key: pseudonymised data remain personal data.

Controller (Art. 4(7)) — entity determining purposes and means. "Within the same controller" = re-identification key may stay internal.

...

Responsibility: (incompleteness)

What counts as sufficient "separately" is not specified.

Who qualifies as "authorised persons" has no exhaustive criteria.

The boundary between pseudonymisation and weak masking is undefined.

Why before analysis? This is like reading definitions and cross-references before interpreting a clause. You load the model's working memory with oriented context — not generic knowledge, but knowledge shaped by your questions.

The model now "sees" the fragment through a prepared lens. Next step: look through it.

STEP 3**Read Through the Question Field**

Combine context (Step 2) with Deep questions as attention geometry and Connect questions as output structure.

PROMPT STRUCTURE**System message:****Attention geometry:**

[Deep questions from Step 1]

Projection structure:

[Connect questions from Step 1]

Output format: Truth .. / Deep .. / Connect .. / Service .. / Knowledge .. / Evolution .. / Responsibility ..

User message:**Context:**

[Full output from Step 2]

Fragment:

[Original text of Recital 29]

MODEL OUTPUT (EXAMPLE)**Truth**

— Recital 29 describes a regime of *managed reversibility*, not permanent identity destruction. It encourages, not mandates.

...

Deep

— Pseudonymisation alone is insufficient for compliance. The boundary of "kept separately" is left to risk-based assessment. Re-identification must remain exceptional, not routine.

...

Connect

— Bridges Art. 24 (controller responsibility), Art. 25 (privacy by design), Art. 32 (security). Simultaneously a security measure, a minimisation measure, and a design-level architectural choice.

...

Responsibility

— Open: exact separation standard, authorised person criteria, boundary with anonymisation, practical audit requirements.

...

What happened? Deep questions told the model *what to look for* (boundaries, sufficiency, scope). Connect questions told it *how to organize* (relationships, axes of compliance). The result is structured analysis, not flat summary.

This single step already goes far beyond "summarize." But we're not done.

STEP 4

Crystallize Into Named Fragments

Extract discrete semantic units from the disclosure. Each fragment carries its own question, forces, and connections.

A TYPICAL PROMPT

Extract assertions from the analysis as structured fragments. For each: name, core question, sustaining forces, connections, dependencies, and full text.

A1 Not an obligation, but a strong normative incentive

"What normative force does Recital 29 apply to pseudonymisation given its incentive-based wording?"

Recital 29 expresses regulatory encouragement in favor of pseudonymisation but does not introduce a universal unconditional obligation to always apply it. Depends on: Semantic Center ·

Connects to: A7, A12

A6 "Kept separately" means real separation, not nominal

"What degree of separation does additional information require given risk and operational accessibility?"

If the same people in the same workflow can easily link code to identity, the "kept separately" requirement is not substantively met. Depends on: A5 · Connects to: A4, A11, A15

A8 Controlled reversibility preserves personal data status

"What legal status does reversibly pseudonymised data produce given additional information can restore identity?"

If re-attribution remains possible via separately stored additional information, data typically remain personal data in pseudonymised form — not anonymous data. Depends on: A5, A6 · Connects to: A12

Why this matters: You now have *named, addressable, connected* semantic units — not a wall of text. Each fragment can be versioned, challenged, reused, or traced back to the questions that produced it.

Summary vs. Semantic Map

SINGLE-PASS SUMMARY

Recital 29 encourages pseudonymisation. Controllers should keep re-identification keys separate and designate authorised persons.


- ✗ No definitions anchored
- ✗ No boundaries identified
- ✗ No connections mapped
- ✗ No open questions surfaced
- ✗ Looks complete — isn't

SEMANTIC KERNEL OUTPUT

- named assertions with dependency graph.
- definitions anchored to source articles.
- boundary questions with explicit scope.
- connection vectors to surrounding provisions.
- open zones explicitly marked as unresolved.

- ✓ Navigable structure
- ✓ Traceable to questions
- ✓ Honest about incompleteness
- ✓ Reusable across analyses

The Protocol at a Glance



Each step uses a different part of the question spectrum. Truth questions load context. Deep questions shape attention. Connect questions organize output. The result is a multi-layered semantic map, not a flat response.

Why questions?

Questions are structurally more stable than answers. A question oriented toward boundaries *returns boundary information*, not arbitrary content. Questions act not only as queries, but as stable semantic anchors — pointers that define regions of meaning even before specific answers are formed.

What to Do with the Fragments

Branch inside your chat

Both ChatGPT and Claude let you branch conversations. Use this to create parallel analysis tracks:

WORKFLOW EXAMPLE

Main thread: Run the protocol on Recital 29

↳ Branch A:

Deep-dive on assertion A6 ("kept separately" – what does it mean in practice?)

↳ Branch B:

Run the protocol on Recital 26

(anonymisation boundary – feeds back into A8) ↳ Branch C:

Cross-reference A1 with Article 25 (is the "incentive" reading defensible?)

Build cumulative context

When analyzing multiple fragments, copy key outputs into the next analysis as pre-loaded context:

NEXT ANALYSIS — PRE-LOADED CONTEXT

Previously established from Recital 29:

- Pseudonymised data remain personal data (A8) - "Kept separately" requires real, not nominal separation (A6) - Pseudonymisation is simultaneously security, minimisation, and privacy-by-design measure (A12)

New fragment:

GDPR Article 25 – Data protection by design and by default

...

Each analysis enriches the field for the next. Fragments accumulate into a **portable semantic toolkit** for your domain.

Fragment types you'll produce

- Assertions — what the text claims (with traceable foundations)
- Definitions — what terms mean *in this specific context*
- Boundaries — where the fragment's applicability starts and stops
- Connections — how it links to surrounding provisions
- Open questions — what it genuinely leaves unresolved

Beyond Chat: What Changes with Automation

What works well in chat mode

- No infrastructure needed — start immediately in any LLM interface
- Teaches the *feel* of question-driven reading through practice
- Produces useful artifacts for individual regulatory analysis

What chat mode can't do

- No persistent state — fragments don't survive between sessions
- No real retrieval — context relies on model knowledge, not actual documents
- No automated tracing — verification requires manual effort
- No cross-fragment coherence — context windows limit accumulation

Automation

The protocol you've just walked through has a property that matters: every step has a defined input and a defined output. Questions go in, structured context comes out. Context goes in, disclosed fragments come out. This means the protocol doesn't just work in chat — it is already shaped for automation.

PIPELINE

Input Parser (fragment metadata)

- **Question Generator** (orientation prompt + LLM call)
- **Context Builder** (Truth Qs + retrieval from Qdrant / doc store)
- **Disclosure Engine** (Deep + Connect Qs + assembled context → LLM)
- **Fragment Extractor** (structured parser → JSON/YAML)
- **Knowledge Store** (persistent graph + vector index)

Automation adds what chat cannot: **persistent fragments** across sessions, **real retrieval** from document stores, **automated completeness checking**, and the ability to run the protocol across an entire regulation systematically.

Upcoming publications will show the protocol running as an automated pipeline — with real retrieval, real fragment accumulation, and real cross-document coherence.

The chat protocol is not a degraded version. It's the seed — prove the structure works at the scale of one person and one fragment, then scale.

*The most valuable thing you can give
someone is not an answer.*

It is a better way of asking.

Semantic Kernel Prompting

A practitioner's guide from the CoAgnes Framework

Full worked example with complete prompts,
step formalization table, and GDPR Recital 29
protocol trace available separately.

CoAgnes Framework · March 2026